

# In-Class Exercise: Power Analysis

XDASI Fall 2021

11/4/2021

## Contents

Example . . . . .	1
Manual power calculation . . . . .	1
Compute power in R . . . . .	2
What if you change different variables that influence power? . . . . .	3
Design for a targeted power . . . . .	5

## Example

**Aho, Ex. 6.10 (also see Fig. 6.7)**

We will set up this example and go through it together in class.

### Manual power calculation

First, compute the power by hand:

```
# ===== #
# set up variables

effect.size = -7 # effect size =  $Exp(X)$  under  $H_A$ 
n = 200
sigma = 45
alpha = 0.05
type = "one.sample" # one or two sample
alt = "one.sided" # one- or two-sided

# set  $Exp(X) = 0$  under null  $H_o$ 

# critical value ( $z^*$ ) for lower-tail test at  $\alpha=0.05$ 
z.crit = qnorm(alpha) # -1.644854

# check alpha using standard normal distribution
pnorm(0,abs(z.crit),lower.tail=T) # 0.05
## [1] 0.05

# compute SEM for sample size
sem = sigma/sqrt(n)
```

```
sem
## [1] 3.181981

# ===== #
# manual power calculation

# compute power using area under the curve for H_A

# get value of critical x at z.crit for H_A
# want  $P(\bar{X} \leq z.crit * sem)$ 
# percent difference for lower-tail significance
x.crit = z.crit*sem # x-value at critical z-score
x.crit
## [1] -5.233892

# Expected  $\bar{X}$  (pop. mean) under H_A is  $(\mu_o - \mu_A): Exp(X) = -7$ 
pwr = pnorm(x.crit, mean = effect.size, sd = sem) # power = 0.71
pwr
## [1] 0.7105643

# check z-score for H_A at expected power
qnorm(pwr, effect.size, sem) # alpha =  $P(\bar{X} \leq -5.24)$ 
## [1] -5.233892
```

## Compute power in R

Given any 4 of the 5 variables that go into the power equation, we can use `power.t.test()` to compute the missing value. Since  $n$  is large, we could also use the `power.z.test()` command from the `asbio` package. These give slightly different results, as the  $t$ -test is a bit more conservative. (They also use different names for their arguments, and the objects they produce are also different.)

**NOTE:** Effect size used for these functions should be given as a positive number, otherwise these functions will not work as expected.

```
# ===== #
# provide expected effect size as a positive number
power.t.test(n, delta = abs(effect.size), sd = sigma, sig.level = alpha,
             type="one.sample", alternative="one.sided", strict=T)

##
##      One-sample t test power calculation
##
##              n = 200
##             delta = 7
##              sd = 45
##      sig.level = 0.05
##              power = 0.7079982
##      alternative = one.sided

# note that arguments for this command differ
power.z.test(n, effect = abs(effect.size), sigma = sigma,
             alpha = alpha, test="one.tail", strict=T)
```

```
## $sigma
## [1] 45
##
## $n
## [1] 200
##
## $power
## [1] 0.7105643
##
## $alpha
## [1] 0.05
##
## $effect
## [1] 7
##
## $test
## [1] "one.tail"
```

What if you change different variables that influence power?

- Increase effect size => increase power (reduce Type II error)
- Increase sample size => increase power (reduce Type II error)
- Raise  $\alpha$  => lower stringency (increase Type I error)

We can compute the new power by hand, or use the `power.z.test()` command:

```
# ===== #
# increase effect size
# ===== #
# what happens if E = -8? => increase power
# (keep alpha the same)
pnorm(x.crit, effect.size - 1, sem) # power = 0.81
```

```
## [1] 0.8076595
```

```
power.z.test(n, effect = -(effect.size-1), sigma = sigma,
             alpha = alpha, test="one.tail", strict=T)
```

```
## $sigma
## [1] 45
##
## $n
## [1] 200
##
## $power
## [1] 0.8076595
##
## $alpha
## [1] 0.05
##
## $effect
## [1] 8
```

```
##
## $test
## [1] "one.tail"
```

```
# ===== #
# increase sample size
# ===== #
# what if sample size = 300? => more power for same E
# (keep alpha the same)
n = 300

# get x-bar and SEM
sem = sigma / sqrt(n)
sem
```

```
## [1] 2.598076
```

```
x.crit = qnorm(0.05)*sem
x.crit
```

```
## [1] -4.273455
```

```
# power
pnorm(x.crit, effect.size, sem) # power = 0.85
```

```
## [1] 0.8530139
```

```
power.z.test(n, effect = -effect.size, sigma = sigma,
             alpha = alpha, test="one.tail", strict=T)
```

```
## $sigma
## [1] 45
##
## $n
## [1] 300
##
## $power
## [1] 0.8530139
##
## $alpha
## [1] 0.05
##
## $effect
## [1] 7
##
## $test
## [1] "one.tail"
```

```
# ===== #
# relax stringency: raise alpha
# ===== #
```

```

# raising alpha increases Type I error
# what happens to power? => power goes down
alpha = 0.2
z.crit = qnorm(alpha) # -0.842

# check alpha2 using standard normal distribution
pnorm(0,abs(z.crit),lower.tail=T)

## [1] 0.2

x.crit = z.crit*sem
x.crit

## [1] -2.186596

pwr = pnorm(x.crit, mean = effect.size, sd = sem) # power = 0.913
pwr

## [1] 0.9680359

qnorm(pwr, effect.size, sem) # check power

## [1] -2.186596

# power is the same with the z-test power function
power.z.test(n, effect = -effect.size, sigma = sigma,
             alpha = alpha, test="one.tail", strict=T)

## $sigma
## [1] 45
##
## $n
## [1] 300
##
## $power
## [1] 0.9680359
##
## $alpha
## [1] 0.2
##
## $effect
## [1] 7
##
## $test
## [1] "one.tail"

```

### Design for a targeted power

What if you want to design the experiment for power = 0.8, for the same sample size and effect size? What is the Type II error? What happens to the Type I error?

```
# ===== #
# increase desired power to 0.8
# ===== #
# if raise desired power without increasing effect size,
# => alpha goes up (less stringent)
x.bar2 = qnorm(0.8, effect.size, sem) # effect size -4.32
x.bar2
```

```
## [1] -4.813404
```

```
# what significance level is this?
alpha2 = x.bar2 / sem
alpha2
```

```
## [1] -1.85268
```

```
pnorm(0,abs(alpha2),lower.tail=T) # 0.087
```

```
## [1] 0.03196412
```

```
# what significance level is this?
alpha2 = x.bar2 / sem
alpha2
```

```
## [1] -1.85268
```

```
pnorm(0,abs(alpha2),lower.tail=T) # 0.087
```

```
## [1] 0.03196412
```

```
# ===== #
# using power.t.test command
# now supply power and ask what new alpha is
power.t.test(n, delta = -effect.size, sd = sigma,
             sig.level = NULL, power = 0.8,
             type="one.sample", alternative="one.sided", strict=T)
```

```
## Warning in pt(qt(sig.level/tside, nu, lower.tail = FALSE), nu, ncp = sqrt(n/
## tsample) * : full precision may not have been achieved in 'pnt{final}'
```

```
##
##      One-sample t test power calculation
##
##              n = 300
##             delta = 7
##              sd = 45
##      sig.level = 0.03251671
##              power = 0.8
##      alternative = one.sided
```

*# gives  $\alpha = 0.088$*