

In-Class Exercise: Power Analysis

XDASI Fall 2021

11/8/2021

Contents

Example	1
What is the question?	1
Manual power calculation	2
Compute power in R	3
What if you change different variables that influence power?	4
Design for a targeted power	6

Example

Aho, Ex. 6.10 (also see Fig. 6.7)

Surprisingly, the incidence of Alzheimer's disease is negatively associated with moderate smoking, possibly because nicotine may reduce apoptosis (programmed cell death) of neurons. The description of the problem is as follows:

EXAMPLE 6.10 POWER ANALYSIS, SMOKING, AND ALZHEIMER'S

Surprisingly, the incidence of Alzheimer's disease has been negatively associated with moderate amounts of smoking (Van Duijn and Hoffman 1991, Graves et al. 1991, Brenner et al. 1993, Salib and Hillier 1997). Smoking may offer some protection because nicotine may reduce apoptosis (programmed cell death) of neurons (Larrick 1993).

Let us assume that previous to the start of the experiment concerning the effect of smoking on Alzheimer's, researchers were interested in being able to detect a 7% decrease in Alzheimer's incidence for subjects smoking 10–20 cigarettes a day, given that $\sigma = 45\%$. The investigators wanted to know if a sample size of 200 was sufficient to detect this effect (produce a significant result) given $\alpha = 0.05$.

Because we know σ , we will run the power analysis assuming that a one-sample z-test will be used for hypothesis testing. We are interested in the lower-tailed alternative hypothesis that smoking decreases the occurrence of Alzheimer's. Our hypotheses are $H_0: \mu \geq 0$ and $H_A: \mu < 0$. We note that we have a sample size large enough to assume a normal sampling distribution for \bar{X} , regardless of its underlying parent distribution.

A one-sample lower-tailed z-test would reject null at $\alpha = 0.05$ whenever z^* is less than the lower-tailed critical value -1.645 .

What is the question?

Researchers were interested in being able to detect a 7% reduction in Alzheimer's for patients smoking 10–20 cigarettes a day, given that $\sigma = 45\%$.

Is a sample size of 200 sufficient to detect this effect size given $\alpha = 0.05$?

- What are H_0 and H_A ?

- Can we assume that the sampling distribution of the sample mean, \bar{X} , is normal?
- What type of test should we use?
- What is the critical value for the test?

your answer here

Manual power calculation

First, compute the power by hand:

```
# ===== #
# set up variables

effect.size = -7 # effect size = Exp(X) under H_A
n = 200
sigma = 45
alpha = 0.05
type = "one.sample" # one or two sample
alt = "one.sided" # one- or two-sided

# set Exp(X) = 0 under null H_o

# critical value (z*) for lower-tail test at alpha=0.05
z.crit = qnorm(alpha) # -1.644854

# check alpha using standard normal distribution
pnorm(0,abs(z.crit),lower.tail=T) # 0.05
## [1] 0.05

# compute SEM for sample size
sem = sigma/sqrt(n)
sem
## [1] 3.181981

# ===== #
# manual power calculation

# compute power using area under the curve for H_A

# get value of critical x at z.crit for H_A
# want P( X.bar <= z.crit * sem)
# percent difference for lower-tail significance
x.crit = z.crit*sem # x-value at critical z-score
x.crit
## [1] -5.233892

# Expected X.bar (pop. mean) under H_A is (mu_o - mu_A): Exp(X) = -7
pwr = pnorm(x.crit, mean = effect.size, sd = sem) # power = 0.71
pwr
## [1] 0.7105643

# check z-score for H_A at expected power
```

```
qnorm(pwr, effect.size, sem) # alpha = P(X.bar <= -5.24)
## [1] -5.233892
```

Compute power in R

Given any 4 of the 5 variables that go into the power equation, we can use `power.t.test()` to compute the missing value. Since n is large, we could also use the `power.z.test()` command from the `asbio` package. These give slightly different results, as the t -test is a bit more conservative. (They also use different names for their arguments, and the objects they produce are also different.)

NOTE: *Effect size used for these functions should be given as a positive number, otherwise these functions will not work as expected.*

```
# ===== #
# provide expected effect size as a positive number
power.t.test(n, delta = abs(effect.size), sd = sigma, sig.level = alpha,
             type="one.sample", alternative="one.sided", strict=T)
```

```
##
##      One-sample t test power calculation
##
##              n = 200
##             delta = 7
##              sd = 45
##      sig.level = 0.05
##             power = 0.7079982
##      alternative = one.sided
```

```
# note that arguments for this command differ
power.z.test(n, effect = abs(effect.size), sigma = sigma,
             alpha = alpha, test="one.tail", strict=T)
```

```
## $sigma
## [1] 45
##
## $n
## [1] 200
##
## $power
## [1] 0.7105643
##
## $alpha
## [1] 0.05
##
## $effect
## [1] 7
##
## $test
## [1] "one.tail"
```

What if you change different variables that influence power?

- Increase effect size => increase power (reduce Type II error)
- Increase sample size => increase power (reduce Type II error)
- Raise α => lower stringency (increase Type I error)

We can compute the new power by hand, or use the `power.z.test()` command:

```
# ===== #
# increase effect size
# ===== #
# what happens if E = -8? => increase power
# (keep alpha the same)
pnorm(x.crit, effect.size - 1, sem) # power = 0.81
```

```
## [1] 0.8076595
```

```
power.z.test(n, effect = -(effect.size-1), sigma = sigma,
             alpha = alpha, test="one.tail", strict=T)
```

```
## $sigma
## [1] 45
##
## $n
## [1] 200
##
## $power
## [1] 0.8076595
##
## $alpha
## [1] 0.05
##
## $effect
## [1] 8
##
## $test
## [1] "one.tail"
```

```
# ===== #
# increase sample size
# ===== #
# what if sample size = 300? => more power for same E
# (keep alpha the same)
n = 300

# get x-bar and SEM
sem = sigma / sqrt(n)
sem
```

```
## [1] 2.598076
```

```
x.crit = qnorm(0.05)*sem
x.crit
```

```
## [1] -4.273455
```

```
# power
pnorm(x.crit, effect.size, sem) # power = 0.85
```

```
## [1] 0.8530139
```

```
power.z.test(n, effect = -effect.size, sigma = sigma,
             alpha = alpha, test="one.tail", strict=T)
```

```
## $sigma
## [1] 45
##
## $n
## [1] 300
##
## $power
## [1] 0.8530139
##
## $alpha
## [1] 0.05
##
## $effect
## [1] 7
##
## $test
## [1] "one.tail"
```

```
# ===== #
# relax stringency: raise alpha
# ===== #
# raising alpha increases Type I error
# what happens to power? => power goes down
alpha = 0.2
z.crit = qnorm(alpha) # -0.842

# check alpha2 using standard normal distribution
pnorm(0,abs(z.crit),lower.tail=T)
```

```
## [1] 0.2
```

```
x.crit = z.crit*sem
x.crit
```

```
## [1] -2.186596
```

```
pwr = pnorm(x.crit, mean = effect.size, sd = sem) # power = 0.913
pwr
```

```
## [1] 0.9680359
```

```
qnorm(pwr, effect.size, sem) # check power
```

```
## [1] -2.186596
```

```
# power is the same with the z-test power function
power.z.test(n, effect = -effect.size, sigma = sigma,
             alpha = alpha, test="one.tail", strict=T)
```

```
## $sigma
## [1] 45
##
## $n
## [1] 300
##
## $power
## [1] 0.9680359
##
## $alpha
## [1] 0.2
##
## $effect
## [1] 7
##
## $test
## [1] "one.tail"
```

Design for a targeted power

What if you want to design the experiment for power = 0.8, for the same sample size and effect size? What is the Type II error? What happens to the Type I error?

```
# ===== #
# increase desired power to 0.8
# ===== #
# if raise desired power without increasing effect size,
# => alpha goes up (less stringent)
x.bar2 = qnorm(0.8, effect.size, sem) # effect size -4.32
x.bar2
```

```
## [1] -4.813404
```

```
# what significance level is this?
alpha2 = x.bar2 / sem
alpha2
```

```
## [1] -1.85268
```

```
pnorm(0,abs(alpha2),lower.tail=T) # 0.087
```

```
## [1] 0.03196412
```

```
# what significance level is this?
```

```
alpha2 = x.bar2 / sem
```

```
alpha2
```

```
## [1] -1.85268
```

```
pnorm(0,abs(alpha2),lower.tail=T) # 0.087
```

```
## [1] 0.03196412
```

```
# ===== #
```

```
# using power.t.test command
```

```
# now supply power and ask what new alpha is
```

```
power.t.test(n, delta = -effect.size, sd = sigma,
```

```
          sig.level = NULL, power = 0.8,
```

```
          type="one.sample", alternative="one.sided", strict=T)
```

```
## Warning in pt(qt(sig.level/tside, nu, lower.tail = FALSE), nu, ncp = sqrt(n/
```

```
## tsample) * : full precision may not have been achieved in 'pnt{final}'
```

```
##
```

```
##       One-sample t test power calculation
```

```
##
```

```
##               n = 300
```

```
##           delta = 7
```

```
##           sd = 45
```

```
##       sig.level = 0.03251671
```

```
##       power = 0.8
```

```
##   alternative = one.sided
```

```
# gives alpha = 0.088
```